

Arrays and Functions

CPSC 275
Introduction to Computer Systems

One-Dimensional Arrays

- An **array** is a data structure containing a number of data values, all of which have the same type.
- These values, known as **elements**, can be individually selected by their position within the array.
- The elements of a one-dimensional array *a* are conceptually arranged one after another in a single row (or column):



2

One-Dimensional Arrays

- To declare an array, we must specify the *type* of the array's elements and the *number* of elements:
- The elements may be of any type; the length of the array can be any (integer) constant expression.
- Using a macro to define the length of an array is an excellent practice:

```
int a[10];

#define N 10
...
int a[N];
```

3

Array Subscripting

- To access an array element, write the array name followed by an integer value in square brackets.
- This is referred to as **subscripting** or **indexing** the array.
- The elements of an array of length *n* are indexed from 0 to *n* - 1.
- If *a* is an array of length 10, its elements are designated by *a* [0], *a* [1], ..., *a* [9]:



4

Array Subscripting

- Expressions of the form *a* [*i*] are lvalues, so they can be used in the same way as ordinary variables:

```
a[0] = 1;
printf("%d\n", a[5]);
++a[i];
```

5

Array Subscripting

- Examples of typical operations on an array *a* of length *N*:

```
for (i = 0; i < N; i++)
    a[i] = 0;           /* clears a */

for (i = 0; i < N; i++)
    scanf("%d", &a[i]); /* reads data into a */

for (i = 0; i < N; i++)
    sum += a[i];        /* sums the elements of a */
```

6

Array Subscripting

- C doesn't require that subscript bounds be checked; if a subscript goes out of range, the program's behavior is undefined.
- A common mistake: forgetting that an array with n elements is indexed from 0 to $n - 1$, not 1 to n :

```
int a[10], i;
for (i = 1; i <= 10; i++)
    a[i] = 0;
```

7

Array Subscripting

- An array subscript may be any integer expression:
- The expression can even have side effects:

```
a[i+j*10] = 0;

i = 0;
while (i < N)
    a[i++] = 0;
```

8

Array Initialization

- An array, like any other variable, can be given an initial value at the time it's declared.
- The most common form of **array initializer** is a list of constant expressions enclosed in braces and separated by commas:

```
int a[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
```

9

Array Initialization

- If the initializer is shorter than the array, the remaining elements of the array are given the value 0:

```
int a[10] = {1, 2, 3, 4, 5, 6};
/* initial value of a is {1, 2, 3, 4, 5, 6, 0, 0, 0, 0} */
```

- Using this feature, we can easily initialize an array to all zeros:

```
int a[10] = {0};
/* initial value of a is {0, 0, 0, 0, 0, 0, 0, 0, 0, 0} */
```

- It's illegal for an initializer to be longer than the array it initializes.

10

Array Initialization

- If an initializer is present, the length of the array may be omitted:

```
int a[] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
```

- The compiler uses the length of the initializer to determine how long the array is.

11

Using the **sizeof** Operator with Arrays

- The **sizeof** operator can determine the size of an array (in bytes).
- If a is an array of 10 integers, then **sizeof(a)** is typically 40 (assuming that each integer requires four bytes).
- We can also use **sizeof** to measure the size of an array element, such as $a[0]$.
- Dividing the array size by the element size gives the length of the array:

```
sizeof(a) / sizeof(a[0])
```

12

Program: Dot Products

Write a C program (**dotprod.c**) which will compute the dot product of two N -vectors, x and y , that is,

$$x \cdot y = x_0 y_0 + \dots + x_{N-1} y_{N-1}$$

Test your program using $x_i = y_i = 1.0, i = 0, \dots, N-1$, for different values of N . Measure the running time of your program when N is large, using Linux **time** command.

Multidimensional Arrays

- An array may have any number of dimensions.
- The following declaration creates a two-dimensional array (a *matrix*, in mathematical terminology):

```
int m[5][9];
```
- m has 5 rows and 9 columns. Both rows and columns are indexed from 0:

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
|---|---|---|---|---|---|---|---|---|---|
| 0 | | | | | | | | | |
| 1 | | | | | | | | | |
| 2 | | | | | | | | | |
| 3 | | | | | | | | | |
| 4 | | | | | | | | | |

14


Multidimensional Arrays

- To access the element of m in row i , column j , we must write $m[i][j]$.
- The expression $m[i]$ designates row i of m , and $m[i][j]$ then selects element j in this row.

15

Multidimensional Arrays

- Although we visualize two-dimensional arrays as tables, that's not the way they're actually stored in computer memory.
- C stores arrays in **row-major order**, with row 0 first, then row 1, and so forth.
- How the m array is stored:



16

Multidimensional Arrays

- Nested `for` loops are ideal for processing multidimensional arrays.
- Consider the problem of initializing an array for use as an *identity matrix*. A pair of nested `for` loops is perfect:

```
#define N 10
double ident[N][N];
int row, col;
for (row = 0; row < N; row++)
    for (col = 0; col < N; col++)
        if (row == col)
            ident[row][col] = 1.0;
        else
            ident[row][col] = 0.0;
```

17

Initializing a Multidimensional Array

- We can create an initializer for a two-dimensional array by nesting one-dimensional initializers:

```
int m[5][9] = {{1, 1, 1, 1, 1, 0, 1, 1, 1},
               {0, 1, 0, 1, 0, 1, 0, 1, 0},
               {0, 1, 0, 1, 1, 0, 0, 1, 0},
               {1, 1, 0, 1, 0, 0, 0, 1, 0},
               {1, 1, 0, 1, 0, 0, 1, 1, 1}};
```

- Initializers for higher-dimensional arrays are constructed in a similar fashion.

18

Program: Matrix-Vector Multiplication

Write a C program (**matvecmul.c**) which will compute the product of an N -by- N matrix A by an N -vector x . Test your program using

$$\begin{aligned} a_{ij} &= i+1, & \text{for } i, j = 0, \dots, N-1 \\ x_j &= 1.0 & \text{for } j = 0, \dots, N-1 \end{aligned}$$

for different values of N . Measure the running time of your program when N is large, using Linux **time** command.

Functions

Function Definitions

- General form of a **function definition**:

```
return-type function-name ( parameters )
{
    declarations
    statements
}
```

- Functions may not return arrays.
- Specifying that the return type is `void` indicates that the function doesn't return a value.

21

Function Definitions

```
#include <stdio.h>

double average(double a, double b) /* function definition */
{
    return (a + b) / 2;
}

int main(void)
{
    double x, y, z;

    printf("Enter three numbers: ");
    scanf("%lf%lf%lf", &x, &y, &z);
    printf("Average of %g and %g: %g\n", x, y, average(x, y));
    printf("Average of %g and %g: %g\n", y, z, average(y, z));
    printf("Average of %g and %g: %g\n", x, z, average(x, z));

    return 0;
}
```

22

Function Declarations

- A **function declaration** provides the compiler with a brief glimpse at a function whose full definition will appear later.
- General form of a function declaration:
`return-type function-name (parameters) ;`
- The declaration of a function must be consistent with the function's definition.
- Here's the `average.c` program with a declaration of `average` added.

23

Function Declarations

```
#include <stdio.h>

double average(double a, double b); /* DECLARATION */

int main(void)
{
    double x, y, z;

    printf("Enter three numbers: ");
    scanf("%lf%lf%lf", &x, &y, &z);
    printf("Average of %g and %g: %g\n", x, y, average(x, y));
    printf("Average of %g and %g: %g\n", y, z, average(y, z));
    printf("Average of %g and %g: %g\n", x, z, average(x, z));

    return 0;
}

double average(double a, double b) /* DEFINITION */
{
    return (a + b) / 2;
}
```

24

Function Declarations

- Function declarations of the kind we're discussing are known as **function prototypes**.
- A function prototype doesn't have to specify the names of the function's parameters, as long as their types are present:

```
double average(double, double);
```

25

Arguments

- In C, arguments are **passed by value**: when a function is called, each argument is evaluated and its value assigned to the corresponding parameter.
- Since the parameter contains a copy of the argument's value, any changes made to the parameter during the execution of the function don't affect the argument.

26

Array Arguments

- When a function parameter is a one-dimensional array, the length of the array can be left unspecified:

```
int f(int a[]) /* no length specified */
{
    ...
}
```
- C doesn't provide any easy way for a function to determine the length of an array passed to it.
- Instead, we'll have to supply the length—if the function needs it—as an additional argument.

27

Array Arguments

- Example:

```
int sum_array(int a[], int n)
{
    int i, sum = 0;

    for (i = 0; i < n; i++)
        sum += a[i];

    return sum;
}
```
- Since `sum_array` needs to know the length of `a`, we must supply it as a second argument.

28

Array Arguments

- The prototype for `sum_array` has the following appearance:

```
int sum_array(int a[], int n);
```
- As usual, we can omit the parameter names if we wish:

```
int sum_array(int [], int);
```

29

Array Arguments

- When `sum_array` is called, the first argument will be the name of an array, and the second will be its length:

```
#define LEN 100

int main(void)
{
    int b[LEN], total;
    ...
    total = sum_array(b, LEN);
    ...
}
```

30

Array Arguments, cont'd

- Notice that we don't put brackets after an array name when passing it to a function:

```
total = sum_array(b[], LEN);  /** WRONG **/  
total = sum_array(b, LEN);   /** CORRECT **/  
      ↑
```

But what are we really passing?

The address of the first element of the array.

```
total = sum_array(&b[0], LEN);
```

31

Passing Multidimensional Arrays

```
#define LEN 100  
  
int main(void)  
{  
    int b[LEN][LEN];  
    ...  
    setToZero(b, LEN);  
    ...  
}  
  
void setToZero(int b[LEN][LEN], int n)    OR  
{ ... }  
  
void setToZero(int b[][LEN], int n)  
{ ... }
```

32

The **return** Statement

- A non-void function must use the **return** statement to specify what value it will return.
- The **return** statement has the form
`return expression ;`
- The expression is often just a constant or variable:

```
return 0;  
return status;
```

- More complex expressions are possible:

```
return n >= 0 ? n : 0;
```

33

dotprod1.c

Rewrite **dotprod.c** which will compute the dot product of two N -vectors using a function. Your **main()** function should define the two vectors, x and y , and call a function which will accept two vectors and their length and return their dot product.

matvecmul1.c

Rewrite **matvecmul.c** which will compute the product of an N -by- N matrix A by an N -vector x using a function, that is, the function should place the product in a vector b , where $b = Ax$.

Your **main()** function should define A , b , and x and call a function which will accept a matrix A , two vectors, x and y , and their length and place the product in b .