


# Data Alignment

CPSC 275  
Introduction to Computer Systems

## Structures & Alignment


### Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

### Aligned Data

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$



## Alignment Principles

- Aligned Data**
  - Primitive data type requires  $K$  bytes
  - Address must be multiple of  $K$
  - Required on some machines; advised on IA32
    - Treated differently by IA32 Linux, x86-64 Linux, and Windows!
- Motivation for Aligning Data**
  - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
    - Inefficient to load or store datum that spans quad word boundaries
    - Virtual memory very tricky when datum spans 2 pages (TBD later)
- Compiler**
  - Inserts gaps in structure to ensure correct alignment of fields, e.g., `.align` directive

## Specific Cases of Alignment (IA32)

- 1 byte: `char`, ...**
  - no restrictions on address
- 2 bytes: `short`, ...**
  - lowest 1 bit of address must be 0
- 4 bytes: `int`, `float`, `char *`, ...**
  - lowest 2 bits of address must be 00
- 8 bytes: `double`, ...**
  - Windows (and most other OS's & instruction sets):
    - lowest 3 bits of address must be 000
  - Linux:
    - lowest 2 bits of address must be 00
    - i.e., treated the same as a 4-byte primitive data type
- 12 bytes: `long double`**
  - Windows, Linux:
    - lowest 2 bits of address must be 00
    - i.e., treated the same as a 4-byte primitive data type


## Specific Cases of Alignment (x86-64)

- 1 byte: `char`, ...**
  - no restrictions on address
- 2 bytes: `short`, ...**
  - lowest 1 bit of address must be 0
- 4 bytes: `int`, `float`, ...**
  - lowest 2 bits of address must be 00
- 8 bytes: `double`, `char *`, ...**
  - Windows & Linux:
    - lowest 3 bits of address must be 000
- 16 bytes: `long double`**
  - Linux:
    - lowest 3 bits of address must be 000
    - i.e., treated the same as a 8-byte primitive data type

## Satisfying Alignment with Structures

- Within structure:**
  - Must satisfy each element's alignment requirement
- Overall structure placement**
  - Each structure has alignment requirement  $K$ 
    - $K$  = Largest alignment of any element
  - Initial address & structure length must be multiples of  $K$
- Example (under Windows or x86-64):**
  - $K = 8$ , due to `double` element


```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```




## Different Alignment Conventions

- x86-64 or IA32 Windows:
  - $K = 8$ , due to **double** element

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```




- IA32 Linux
  - $K = 4$ ; **double** treated like a 4-byte data type



## Meeting Overall Alignment Requirement

- For largest alignment requirement  $K$
- Overall structure must be multiple of  $K$


```
struct S2 {
    double v;
    int i[2];
    char c;
} *p;
```



## Arrays of Structures

- Overall structure length multiple of  $K$
- Satisfy alignment requirement for every element


```
struct S2 {
    double v;
    int i[2];
    char c;
} a[10];
```



## Accessing Array Elements

- Element  $j$  is at offset 8 within structure
- Assembler gives offset  $a+8$
- Compute array offset  $12i$  for element  $i$ 
  - $\text{sizeof}(S3)$ , including alignment spacers

```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```



## Saving Space

- Put large data types first

```
struct S4 {
    char c;
    int i;
    char d;
} *p;
```



```
struct S5 {
    int i;
    char c;
    char d;
} *p;
```


- Effect ( $K=4$ )




## Union Allocation

- Allocate according to largest element
- Can only use one field at a time

```
union U1 {
    char c;
    int i[2];
    double v;
} *up;
```



```
struct S1 {
    char c;
    int i[2];
    double v;
} *sp;
```



## Byte Ordering, Revisited

### Idea

- Short/long/quad words stored in memory as 2/4/8 consecutive bytes
- Which is most (least) significant?
- Can cause problems when exchanging binary data between machines

### Big Endian

- Most significant byte has lowest address
- Sun Sparc

### Little Endian

- Least significant byte has lowest address
- Intel x86

## Byte Ordering Example

```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```

32-bit

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

64-bit

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

## Byte Ordering Example (Cont).

```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 ==\n",
       [0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",
       dw.c[0], dw.c[1], dw.c[2], dw.c[3],
       dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",
       dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x]\n", dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx]\n", dw.l[0]);
```

## Byte Ordering on IA32

Little Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

LSB ← Print → MSB

Output:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
 Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]  
 Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]  
 Long 0 == [0xf3f2f1f0]

## Byte Ordering on Sun

Big Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

MSB ← Print → LSB

Output on Sun:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
 Shorts 0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]  
 Ints 0-1 == [0xf0f1f2f3,0xf4f5f6f7]  
 Long 0 == [0xf0f1f2f3]

## Byte Ordering on x86-64

Little Endian

f0	f1	f2	f3	f4	f5	f6	f7
c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

LSB ← Print → MSB

Output on x86-64:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]  
 Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]  
 Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]  
 Long 0 == [0xf7f6f5f4f3f2f1f0]

## Practice Problems

- Read CSaPP Sec. 3.9.2 and 3.9.3 and try the following problems:  
3.41 and 3.42