

Processes


CPSC 275
Introduction to Computer Systems

Processes

- Definition: A *process* is an instance of a running program.
 - One of the most profound ideas in computer science
 - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
 - Logical control flow
 - Each program seems to have exclusive use of the CPU
 - Private virtual address space
 - Each program seems to have exclusive use of main memory
- How are these illusions maintained?
 - Process executions interleaved (multitasking) or run on separate cores
 - Address spaces managed by virtual memory system


Memory Layout of a Process

- Text segment**
 - machine-language instructions
 - Read-only
- Initialized data segment**
 - global and static that are explicitly initialized
- Uninitialized data segment**
 - global and static that are not explicitly initialized
 - Initialized to 0 when starting the program
- Stack:**
 - Stores stack frames, one for each function called
- Heap**
 - Dynamically allocated memory



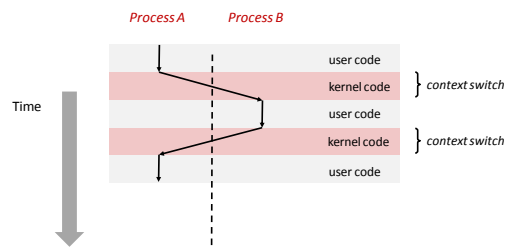
Concurrent Processes

- Two processes *run concurrently* if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



Context Switching

- Control flow passes from one process to another via a *context switch*



Creating New Processes

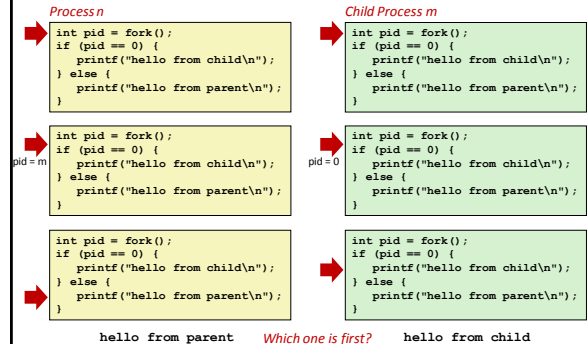
int fork(void)

- Creates a new process (*child* process) that is identical to the calling process (*parent* process)

```
int pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

- Called *once* but returns *twice*
 - returns 0 to the child process
 - returns child's **pid** to the parent process

Understanding fork



fork Example #1

- Parent and child both run same code
 - Distinguish parent from child by return value from **fork**
- Start with same state, but each has private copy
 - Including shared output file descriptor
 - Relative ordering of their print statements undefined

```
void fork1()
{
    int x = 1;
    int pid = fork();
    if (pid == 0) {
        printf("Child has x = %d\n", ++x);
    } else {
        printf("Parent has x = %d\n", --x);
    }
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```

fork Example #2

- Both parent and child can continue


```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



fork Example #3

- Both parent and child can continue


```
void fork3()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("L2\n");
    fork();
    printf("Bye\n");
}
```



fork Example #4

- Both parent and child can continue


```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



fork Example #5

- Both parent and child can continue

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
            fork();
        }
    }
    printf("Bye\n");
}
```



Ending a process

void exit(int status)

- exits a process
 - Normally return with status 0
- atexit()** registers functions to be executed upon exit

```
void cleanup(void) {
    printf("cleaning up\n");
}

void fork6() {
    atexit(cleanup);
    fork();
    exit(0);
}
```

Zombies

- Idea
 - When process terminates, still consumes system resources
 - Various tables maintained by OS
 - Called a "zombie"
- Reaping
 - Performed by parent on terminated child
 - Parent is given exit status information
 - Kernel discards process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then child will be reaped by **init** process

Zombie Example

```
void fork7()
{
    if (fork() == 0) { /* child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else { /* parent */
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* infinite loop */
    }
}
```

Synchronizing with Children


int wait(int *child_status)

- suspends current process until one of its children terminates
- return value is the **pid** of the child process that terminated
- if **child_status != NULL**, then the object it points to will be set to a status indicating why the child process terminated

Example: Synchronizing with Children

```
void fork9() {
    int child_status;

    if (fork() == 0) {
        printf("HC: hello from child\n");
    }
    else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
    exit();
}
```



Checking Exit Status of Children

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10()
{
    int pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* child */
    for (i = 0; i < N; i++) {
        int wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```

Waiting for a Specific Process

waitpid(pid, &status, options)

- suspends current process until specific process terminates
- various options (see manpage of waitpid())

```
void fork11()
{
    int pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = N-1; i >= 0; i--) {
        int wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

Other Useful Functions on Processes

sleep(n)

- suspends current process for **n** seconds.

exec()

- a family of functions to load and run a new program in the context of the current process.
- Lab 11 on Monday