


Virtual Memory


CPSC 275
Introduction to Computer Systems

A System Using Physical Addressing



- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

A System Using Virtual Addressing



- Used in all modern computer systems.

Address Spaces


- Virtual address space:** Set of $N = 2^n$ virtual addresses $\{0, 1, 2, 3, \dots, N-1\}$
- Physical address space:** Set of $M = 2^m$ physical addresses $\{0, 1, 2, 3, \dots, M-1\}$
- Every byte in main memory:
 - one physical address
 - one (or more) virtual addresses

Why Virtual Memory (VM)?

- Uses main memory efficiently
 - Use DRAM as a cache for the parts of a virtual address space
- Simplifies memory management
 - Each process gets the same uniform linear address space
- Isolates address spaces
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information

VM as a Tool for Caching

- Virtual memory** is an array of N contiguous bytes stored on disk.
- Some contents on disk are cached in **physical memory (DRAM cache)**
 - These cache blocks are called **pages** (size is $P = 2^p$ bytes)

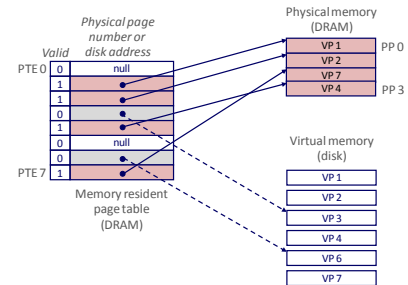


DRAM Cache Organization

- Driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM
 - Disk is about **10,000x** slower than DRAM
- Consequences
 - Large page (block) size: typically 4-8 KB
 - Fully associative:
 - Any VP can be placed in any PP
 - Requires a “large” mapping function
 - Highly sophisticated, expensive replacement algorithms
 - Write-back* rather than *write-through*


Page Tables

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM



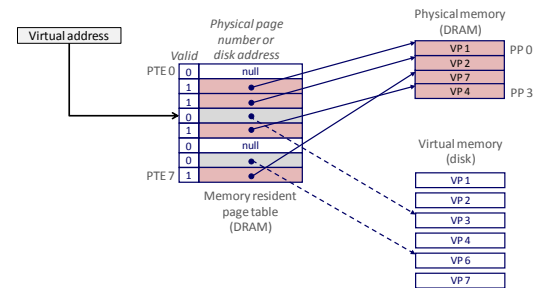
Page Hit

- Page hit:** reference to VM word that is in physical memory




Page Fault

- Page fault:** reference to VM word that is not in physical memory



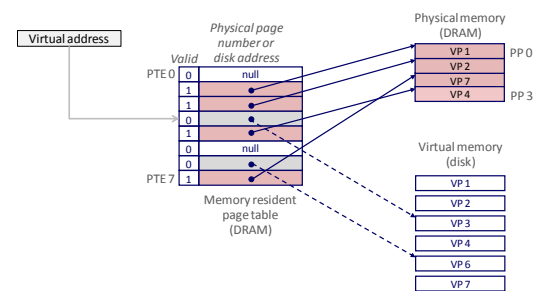
Handling Page Fault

- Page miss causes page fault




Handling Page Fault

- Page miss causes page fault
- Page fault handler selects a victim to be evicted (here VP 4)



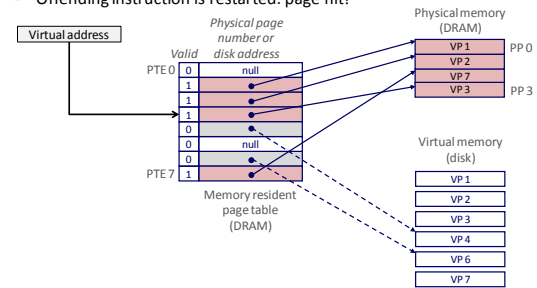
Handling Page Fault

- Page miss causes page fault
- Page fault handler selects a victim to be evicted (here VP 4)



Handling Page Fault

- Page miss causes page fault
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



Locality to the Rescue Again!

- Virtual memory works because of locality
- At any point in time, programs tend to access a set of active virtual pages called the **working set**
 - Programs with better temporal locality will have smaller working sets
- If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- If (SUM(working set sizes) > main memory size)
 - Thrashing**: Performance meltdown where pages are swapped (copied) in and out continuously

VM Address Translation

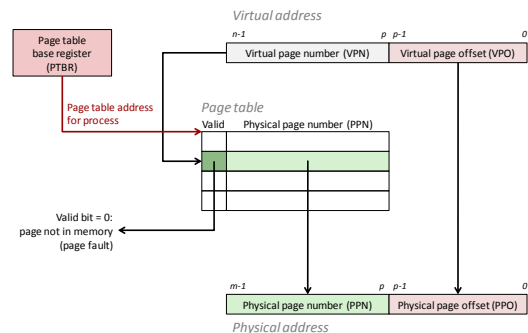
- Virtual Address Space
 $V = \{0, 1, \dots, N-1\}$
- Physical Address Space
 $P = \{0, 1, \dots, M-1\}$
- Address Translation

MAP: $V \rightarrow P \cup \{\emptyset\}$


For virtual address a :

- $\text{MAP}(a) = a'$ if data at virtual address a is at physical address $a' \in P$
- $\text{MAP}(a) = \emptyset$ if data at virtual address a is not in physical memory

Address Translation With a Page Table




Address Translation: Page Hit



- Processor sends virtual address to MMU
- MMU fetches PTE from page table in memory
- MMU sends physical address to cache/memory
- Cache/memory sends data word to processor

Address Translation: Page Fault




- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if *dirty*, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Speeding up Translation with a TLB

Translation Lookaside Buffer (TLB)


- Small hardware cache in MMU
- Maps virtual page numbers to physical page numbers
- Contains complete page table entries for small number of pages

TLB Hit



A TLB hit eliminates a memory access

TLB Miss




A TLB miss incurs an additional memory access

Simple Memory System Example

Addressing

- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



Simple Memory System Page Table

How many entries are in the page table?


Only show first 16 entries (out of 256)

VPN	PPN	Valid
00	28	1
01	–	0
02	33	1
03	02	1
04	–	0
05	16	1
06	–	0
07	–	0

VPN	PPN	Valid
08	13	1
09	17	1
0A	09	1
0B	–	0
0C	–	0
0D	2D	1
0E	11	1
0F	0D	1

Simple Memory System TLB


- 16 entries
- 4-way associative



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	—	0	09	0D	1	00	—	0	07	02	1
1	03	2D	1	02	—	0	04	—	0	0A	—	0
2	02	—	0	08	—	0	06	—	0	03	—	0
3	07	—	0	03	0D	1	0A	34	1	02	—	0

Simple Memory System Cache

- 16 lines, 4-byte block size
- Physically addressed
- Direct mapped




Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	—	—	—	—
2	18	1	00	02	04	08
3	36	0	—	—	—	—
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	—	—	—	—
7	16	1	11	C2	DF	03

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	—	—	—	—
A	2D	1	93	15	DA	38
B	08	0	—	—	—	—
C	12	0	—	—	—	—
D	16	1	04	96	34	15
E	13	1	83	77	1B	03
F	14	0	—	—	—	—


Address Translation Example

Virtual Address: 0x03D4



VPN: 0x0F TLBI: 0x3 TLBT: 0x03 TLB Hit? **Y** Page Fault? **N** PPN: 0x0D

Physical Address



CO: 0 CI: 0x5 CT: 0x0D Hit? **Y** Byte: 0x36

Practice Problems

- Read CSaPP Sec. 9.8-9.9 and try Practice Problems 9.1, 9.2, 9.3, and 9.4.